
Data Documentation Audit

Critique the public documentation for **one dataset or model (or other AI artifact) of your choice** and propose changes that would make it more useful. You may choose any publicly available datasheet, model card, or system card (prof approval is not required).

Work individually. Submit to both Crowdmark (screencaptures of writing for each section) and LEARN Dropbox (pdf of full submission, including any supplemental files) by October 19, 11:59pm.

Submission format

For the documentation artifact you have chosen, write an **executive summary of at most two pages** that analyzes its strengths and weaknesses, and proposes changes that could be made (if you were the model/dataset owner) to improve documentation. Your report may include references and appendices, which do not count against the 2-page limit. Any supplementary files (extra data points, script used to analyze the data) can be submitted to the LEARN dropbox: just zip your submission with the supplements before uploading.

Use the supplied LaTeX template (`a2-review.tex` in `a2-template.zip`), or another tool with equivalent formatting: at least 11-point text and 1-inch margins. Replace the template prompts with your responses. Organize the summary under the three headings below.

Follow the course academic-integrity and AI policies. AI may edit your own prose, with disclosure, but may not generate your analysis or responses. Place any disclosure outside the two-page summary.

[8 marks] Question 1: Context

Identify the dataset or model, its documentation, and a use against which you will assess it. Explain the purpose of the documentation and its intended audience. Consider how the authors describe the scope and intended use case(s) for their dataset or model.

Record the documentation URL, access date, and version or revision where available.

[20 marks] Question 2: Critique

Assess what the documentation communicates well and its important weaknesses. Compare what the documentation offers with relevant recommendations from the literature (e.g. Datasheets (Gebru et al. 2021), Model Cards (Mitchell et al. 2019), or another framework you deem relevant). Explain how these affect a reader's ability to judge suitability for your chosen use. Does the documentation artifact have any notable gaps or limitations? Cite source material (e.g. specific sections or website pages from your chosen artifact) as needed.

[12 marks] Question 3: Proposed changes

Explain how you would change the documentation and why. Prioritize changes that address the findings of your critique. Identify any additional information or evidence you would need, and who could provide it. Example replacement text may be included in an appendix, but is not required. Do not invent missing facts.

Finding documentation

Some places to start:

- **Hugging Face:** browse [datasets](#) and [models](#), then read their dataset or model cards.
- **Kaggle:** browse [datasets](#) and their descriptions, including links to original sources.
- **Industry documentation:** public model and system cards, such as [Anthropic's system cards](#) and [Google DeepMind's model cards](#).
- **Research repositories:** dataset pages, project websites, and documentation linked from research papers. The [UCI Machine Learning Repository](#) is another starting point.

These are suggestions, not a restricted list. Choose documentation with enough substance to support a critical review. Identify who wrote it: a repository's card may differ from the original developer's documentation.

Connecting your analysis to course content

Once you have chosen a dataset, model, or other AI artifact, you may find it helpful to give a close reading to one of the documentation frameworks in the course reading list, such as [Datasheets](#) or [Model Cards](#). Beyond the provided list, many other documentation frameworks have been proposed in the literature, such as [Foundation Model Transparency Reports](#) and [Audit Cards](#), just to name a few. Feel free to choose a framework for your analysis that is suitable for the artifact you have chosen.

Gebru, Timnit, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. 2021. "Datasheets for Datasets." *Communications of the ACM* 64 (12): 86–92. <https://doi.org/10.1145/3458723>.

Mitchell, Margaret, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. 2019. "Model Cards for Model Reporting." In *Proceedings of the Conference on Fairness, Accountability, and Transparency*. <https://doi.org/10.1145/3287560.3287596>.